

Machine Learning Training Module

A Comprehensive Guide for Beginners

Prepared by xAI Training Team

May 07, 2025

Downloadable PDF version available for training purposes.

Contents

1	Introduction to Machine Learning	2
1.1	Why Study Machine Learning?	2
1.2	Setting Up the Environment	2
2	Supervised Learning	2
2.1	Linear Regression Example	2
3	Classification	3
3.1	Logistic Regression Example	3
4	Unsupervised Learning	3
4.1	K-Means Clustering Example	3
5	Model Evaluation	4
5.1	Cross-Validation Example	4
6	Feature Engineering	4
6.1	Standard Scaling Example	5
7	Decision Trees	5
7.1	Decision Tree Classifier Example	5
8	Neural Networks (Introduction)	5
8.1	Simple Neural Network Example	6
9	Hyperparameter Tuning	6
9.1	Grid Search Example	6
10	Conclusion	7
11	References	7

1 Introduction to Machine Learning

Machine Learning (ML) is a field of artificial intelligence that enables systems to learn from data and improve over time. This module introduces ML concepts—supervised and unsupervised learning, model evaluation, and more—with Python examples using scikit-learn.

1.1 Why Study Machine Learning?

- **Automation:** Enables predictive and decision-making systems.
- **Applications:** Used in healthcare, finance, and autonomous systems.
- **Innovation:** Drives advancements in AI and data science.

1.2 Setting Up the Environment

Install Python and libraries like scikit-learn, numpy, and pandas. Use a Jupyter Notebook or IDE like PyCharm. Install dependencies with:

```
1 pip install scikit-learn numpy pandas matplotlib
```

2 Supervised Learning

Supervised learning uses labeled data to predict outcomes.

2.1 Linear Regression Example

Predict a continuous variable using linear regression.

```
1 import numpy as np
2 from sklearn.linear_model import LinearRegression
3 import matplotlib.pyplot as plt
4
5 # Sample data
6 X = np.array([[1], [2], [3], [4], [5]])
7 y = np.array([2, 4, 5, 4, 5])
8
9 # Train model
10 model = LinearRegression()
11 model.fit(X, y)
12
13 # Predict
14 y_pred = model.predict(X)
15
16 # Plot
17 plt.scatter(X, y, color='blue', label='Data')
18 plt.plot(X, y_pred, color='red', label='Fit')
19 plt.xlabel('X')
```

```
20 plt.ylabel('y')
21 plt.legend()
22 plt.savefig('linear_regression.png')
```

Explanation:

- `LinearRegression`: Fits a linear model.
- `model.fit`: Trains the model on input data.
- `plt.savefig`: Saves the plot as an image.

3 Classification

Classification predicts discrete labels.

3.1 Logistic Regression Example

Classify data using logistic regression.

```
1 from sklearn.datasets import load_iris
2 from sklearn.linear_model import LogisticRegression
3 from sklearn.model_selection import train_test_split
4
5 # Load data
6 iris = load_iris()
7 X, y = iris.data, iris.target
8
9 # Split data
10 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size
    =0.2, random_state=42)
11
12 # Train model
13 model = LogisticRegression(max_iter=200)
14 model.fit(X_train, y_train)
15
16 # Evaluate
17 accuracy = model.score(X_test, y_test)
18 print(f"Accuracy: {accuracy:.2f}")
```

4 Unsupervised Learning

Unsupervised learning finds patterns in unlabeled data.

4.1 K-Means Clustering Example

Group data into clusters.

```
1 from sklearn.cluster import KMeans
2 import matplotlib.pyplot as plt
3
4 # Sample data
5 X = np.array([[1, 2], [1.5, 1.8], [5, 8], [8, 8], [1, 0.6], [9,
6     11]])
7
8 # Train model
9 kmeans = KMeans(n_clusters=2, random_state=42)
10 kmeans.fit(X)
11
12 # Plot clusters
13 plt.scatter(X[:, 0], X[:, 1], c=kmeans.labels_, cmap='viridis')
14 plt.scatter(kmeans.cluster_centers_[:, 0], kmeans.cluster_centers_
15    [:, 1], s=200, c='red', marker='X')
16 plt.xlabel('Feature 1')
17 plt.ylabel('Feature 2')
18 plt.savefig('kmeans_clustering.png')
```

5 Model Evaluation

Evaluate models to ensure reliability.

5.1 Cross-Validation Example

Use cross-validation to assess model performance.

```
1 from sklearn.datasets import load_iris
2 from sklearn.model_selection import cross_val_score
3 from sklearn.ensemble import RandomForestClassifier
4
5 # Load data
6 iris = load_iris()
7 X, y = iris.data, iris.target
8
9 # Train model
10 model = RandomForestClassifier(random_state=42)
11 scores = cross_val_score(model, X, y, cv=5)
12
13 # Print results
14 print(f"Cross-validation scores: {scores}")
15 print(f"Average accuracy: {scores.mean():.2f}")
```

6 Feature Engineering

Feature engineering improves model performance by transforming data.

6.1 Standard Scaling Example

Normalize features using StandardScaler.

```
1 from sklearn.preprocessing import StandardScaler
2 import numpy as np
3
4 # Sample data
5 X = np.array([[1, 2], [3, 4], [5, 6], [7, 8]])
6
7 # Scale features
8 scaler = StandardScaler()
9 X_scaled = scaler.fit_transform(X)
10
11 # Print results
12 print("Scaled data:")
13 print(X_scaled)
```

7 Decision Trees

Decision trees make decisions based on feature splits.

7.1 Decision Tree Classifier Example

Classify data using a decision tree.

```
1 from sklearn.datasets import load_iris
2 from sklearn.tree import DecisionTreeClassifier
3 from sklearn.model_selection import train_test_split
4
5 # Load data
6 iris = load_iris()
7 X, y = iris.data, iris.target
8
9 # Split data
10 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size
    =0.2, random_state=42)
11
12 # Train model
13 model = DecisionTreeClassifier(random_state=42)
14 model.fit(X_train, y_train)
15
16 # Evaluate
17 accuracy = model.score(X_test, y_test)
18 print(f"Accuracy: {accuracy:.2f}")
```

8 Neural Networks (Introduction)

Neural networks model complex patterns.

8.1 Simple Neural Network Example

Use scikit-learn's MLPClassifier for a basic neural network.

```
1 from sklearn.neural_network import MLPClassifier
2 from sklearn.datasets import load_iris
3 from sklearn.model_selection import train_test_split
4
5 # Load data
6 iris = load_iris()
7 X, y = iris.data, iris.target
8
9 # Split data
10 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size
    =0.2, random_state=42)
11
12 # Train model
13 model = MLPClassifier(hidden_layer_sizes=(10,), max_iter=1000,
    random_state=42)
14 model.fit(X_train, y_train)
15
16 # Evaluate
17 accuracy = model.score(X_test, y_test)
18 print(f"Accuracy: {accuracy:.2f}")
```

9 Hyperparameter Tuning

Optimize model performance by tuning hyperparameters.

9.1 Grid Search Example

Use GridSearchCV to find the best parameters.

```
1 from sklearn.model_selection import GridSearchCV
2 from sklearn.svm import SVC
3 from sklearn.datasets import load_iris
4
5 # Load data
6 iris = load_iris()
7 X, y = iris.data, iris.target
8
9 # Define parameter grid
10 param_grid = {'C': [0.1, 1, 10], 'kernel': ['linear', 'rbf']}
11
12 # Perform grid search
13 model = GridSearchCV(SVC(), param_grid, cv=5)
14 model.fit(X, y)
15
16 # Print results
17 print(f"Best parameters: {model.best_params_}")
```

```
18 print(f"Best score: {model.best_score_:.2f}")
```

10 Conclusion

This module covers machine learning fundamentals—supervised and unsupervised learning, model evaluation, feature engineering, and neural networks. Practice these examples and explore advanced topics like deep learning with TensorFlow or PyTorch.

11 References

- Scikit-learn Documentation: <https://scikit-learn.org>
- Machine Learning, Tom Mitchell